

Agentic AI for HR

Design human-centred Copilot experiences and governed HR agents across the employee lifecycle.

TGS-2024045795

2 days / 16 hours • Human Resource Digitalisation-4 (HRS-HRM-4031-1.1)



Digital attendance

Complete the attendance check when directed by your trainer.

01 Digital attendance

Scan the trainer-provided QR code at the required checkpoints.

02 General introduction

Course purpose, adult-learning approach and safe practice environment.

03 Dr. Alfred Ang

Lead trainer • AI transformation, analytics and responsible adoption.

Trainer introduction

Learn with practical HR cases and governed Microsoft AI patterns.

01 Digital attendance

Scan the trainer-provided QR code at the required checkpoints.

02 General introduction

Course purpose, adult-learning approach and safe practice environment.

03 Dr. Alfred Ang

Lead trainer • AI transformation, analytics and responsible adoption.

Dr. Alfred Ang

Facilitator for your two-day learning journey.

01 Digital attendance

Scan the trainer-provided QR code at the required checkpoints.

02 General introduction

Course purpose, adult-learning approach and safe practice environment.

03 Dr. Alfred Ang

Lead trainer • AI transformation, analytics and responsible adoption.

Learner introductions

Connect experience to the cases we will solve.

01 Digital attendance

Scan the trainer-provided QR code at the required checkpoints.

02 General introduction

Course purpose, adult-learning approach and safe practice environment.

03 Dr. Alfred Ang

Lead trainer • AI transformation, analytics and responsible adoption.

04 Meet the learning group

Role, HR challenge, Microsoft 365 experience and one question you want answered.

Learning agreements

Build a safe, inclusive and evidence-led learning space.

01 Digital attendance

Scan the trainer-provided QR code at the required checkpoints.

02 General introduction

Course purpose, adult-learning approach and safe practice environment.

03 Dr. Alfred Ang

Lead trainer • AI transformation, analytics and responsible adoption.

Course architecture

Four connected learning units

01

From Generative AI to Agentic HR
Choose the lightest capability that safely solves the work

02

Microsoft 365 Copilot Across the HR Lifecycle
Generate and analyse with a human in control

03

Build Governed HR Agents
Agent Builder for knowledge, Copilot Studio for orchestration

04

Responsible Adoption and HR Transformation
Govern decisions, data, people and agent operations

Learning outcomes

By the end of the course

01 L01

Compare how generative and agentic AI reshape HR operating models versus conventional HR delivery.

02 L02

Analyse where Microsoft 365 Copilot and AI agents can improve end-to-end HR processes and employee experience.

03 L03

Conduct feasibility analysis and design a governed implementation for an HR generative-AI or agentic-AI use case.

04 L04

Apply legal, ethical, privacy, security, fairness and human-oversight controls to AI-enabled HR work.

Technical Skills and Competencies

Human Resource Digitalisation-4 (HRS-HRM-4031-1.1)

01 Knowledge

Generative AI, agentic AI and HR digitalisation

02 Application

Use-case analysis, feasibility, agent design and responsible controls

03 Evidence

Prompts, outputs, checks, decisions and audit trail

Day 1 journey

Frame and generate

01 Morning

HR AI continuum •
operating model •
opportunity
selection

02 Activities 1 –3

Portfolio, workforce
plan and pilot
charter

03 Afternoon

Recruitment, policy,
onboarding with
Microsoft 365
Copilot

04 Activities 4 –6

Evidence matrix,
policy update and
onboarding journey

Day 2 journey

Orchestrate and govern

01 Morning

Agent Builder and
Copilot Studio
patterns

02 Activities 7-9

Policy, onboarding
and leave agents

03 Afternoon

Responsible
adoption and agent
operations

04 Assessment

Written Assessment
+ Case Study
Assessment

Assessment briefing

Two individual instruments • 60 minutes each

01 Written Assessment

Five open-response questions mapped to K1–K5.

02 Case Study Assessment

Five applied questions mapped to A1–A5.

03 Competency standard

Answer every question and meet the stated evidence requirements.

04 Integrity

Use only resources permitted by the assessor; submit your own work.

Assessment flow

Read → plan → answer → verify → submit

01 1. Read

Identify the action verb and required evidence.

02 2. Plan

Allocate time and structure the response.

03 3. Answer

Apply course concepts to the question or case.

04 4. Verify

Check completeness, clarity and source boundaries.

05 5. Submit

Follow assessor instructions and retain confidentiality.

Funding and attendance

Follow the trainer and academy instructions

01 Attendance

Complete required digital check-ins.

02 Assessment

Attempt both assessment instruments.

03 Evidence

Use your own name and candidate details.

04 Support

Raise access or accommodation needs early.

Learning resources

Everything needed for practice

01 Slides

Visual concepts and activity briefs.

02 Learner Guide

Detailed procedures, prompts and reflection.

03 Activity folders

Synthetic cases, workbooks, test data and checklists.

04 GitHub

<https://github.com/tertiarycourses/TGS-2024045795-Agent-ai-for-HR>

Access the activity files

Download or clone before hands-on work

01 Open

<https://github.com/tertiarycourses/TGS-2024045795-Agentic-AI-for-HR>

02 Download

Code → Download ZIP, then extract locally.

03 Locate

Open the numbered activities folder.

04 Protect

Use only the synthetic learner files; never upload live employee data.

From Generative AI to Agentic HR

Choose the lightest capability that safely solves the work

FRAME

GENERATE


ORCHESTRATE

GOVERN

Concepts → evidence → activity → reflection

The HR AI continuum

Automation follows rules; copilots create; agents pursue bounded goals.



Design question • What evidence would make this safe enough for HR work?

Conventional HR operating model

Queues and hand-offs make service slow and opaque.

01

Requests arrive through email and forms

02

Specialists search systems manually

03

Approvals cross organisational boundaries

04

Evidence is scattered across tools

Copilot-assisted HR operating model

A human remains the operator while Copilot accelerates knowledge work.

**Ground prompts in
approved work content**

**A human remains the
operator while Copilot
accelerates knowledge work.**

**Generate a first draft, not
a final decision**

**Compare and refine with
domain expertise**

**Record source and
reviewer**

Agent-enabled HR operating model


An agent can coordinate a bounded workflow across systems.

- 1 Interpret an employee intent**
- 2 Retrieve policy and contextual data**
- 3 Call approved tools and workflows**
- 4 Escalate exceptions with an audit trail**

Human accountability remains visible at every layer.

Assistant, agent or workflow?

Use the simplest lever that meets the outcome and risk profile.



Design question • What evidence would make this safe enough for HR work?

The Microsoft HR AI stack

Microsoft 365 Copilot and Copilot Studio play different but complementary roles.

01

**Word and Outlook: content
and communication**

02

Excel: evidence and analysis

03

**Teams and SharePoint:
collaboration and knowledge**

04

**Copilot Studio: agents, tools,
flows and channels**

Agent anatomy

Reliable agents need a goal, knowledge, instructions, tools, boundaries and tests.

**Goal defines the
employee outcome**

**Reliable agents need a goal,
knowledge, instructions,
tools, boundaries and tests.**

**Knowledge constrains
grounded answers**

**Instructions set
behaviour and refusal
rules**

**Tools act; tests prove
expected behaviour**

HR lifecycle opportunity map


AI value appears across attract, join, grow, serve and exit.

- 1 Attract: workforce planning and recruitment**
- 2 Join: onboarding and access**
- 3 Grow: performance, learning and mobility**
- 4 Serve and exit: policy, leave, cases and offboarding**

Human accountability remains visible at every layer.

Value without losing the human

Saved time should be redirected to judgement, coaching and care.



Design question • What evidence would make this safe enough for HR work?

Use-case selection

Start with frequent, bounded and measurable work.

01

High volume and clear policy

02

Reliable knowledge and integrations

03

Reversible actions and escalation

04

A KPI baseline and accountable owner

Feasibility lens

A good idea is not yet a deployable use case.

**Business value and
employee impact**

**Data readiness and
permissions**

**A good idea is not yet a
deployable use case.**

**Technical reach and
failure modes**

**Risk, change and
operating ownership**

From pilot to operating capability

A pilot must prove usefulness, control and adoption together.

- 1 Define baseline and success threshold**
- 2 Use synthetic data before live data**
- 3 Test normal, edge and adversarial cases**
- 4 Assign monitoring, review and retirement**

Human accountability remains visible at every layer.

Activity 1: Map the HR AI Opportunity Portfolio

Classify HR work across automation, Microsoft 365 Copilot and agentic AI, then select one bounded pilot.

CASE

HarbourLight Logistics SG has 620 employees, a five-person HR team and rising service backlogs. The CHRO wants a defensible starting portfolio—not a list of fashionable tools.

TOOLS

- Microsoft Teams
- Microsoft 365 Copilot
- Excel

OUTPUTS

- Completed opportunity-and-risk matrix
- Two-minute pilot recommendation
- Evidence log

GUARDRAILS

- No employee personal data
- State all assumptions
- Human owner for every shortlisted use case

Detailed procedure, prompts and evidence checklist: Learner Guide + activity folder

Activity 1 success map

Show your reasoning—not only a polished output.

INPUT

**Use the supplied
synthetic case
and approved
sources**

✓ reviewable

DECISION

**State
assumptions,
criteria and the
human owner**

✓ reviewable

EVIDENCE

**Save prompts,
output, checks,
revisions and
approvals**

✓ reviewable

QUALITY

**Pass the
acceptance
criteria and
answer reflection
questions**

✓ reviewable

Stop if the agent or Copilot output cannot be traced, corrected or safely escalated.

Activity 2: Design a Skills-Based Workforce Plan

Use Copilot to synthesise demand evidence and draft a skills-based, fair job design.

CASE

HarbourLight is opening a digitally enabled distribution hub. HR must decide whether to hire, redeploy or develop people for a new Workforce Operations Analyst role.

TOOLS

- Excel with Microsoft 365 Copilot
- Word with Microsoft 365 Copilot

OUTPUTS

- Workforce hypothesis
- Skills gap table
- Inclusive job description

GUARDRAILS

- Separate evidence from assumption
- No age, nationality or other non-job-related criteria
- Hiring manager approval

Detailed procedure, prompts and evidence checklist: Learner Guide + activity folder

Activity 2 success map

Show your reasoning—not only a polished output.

INPUT

**Use the supplied
synthetic case
and approved
sources**

✓ reviewable

DECISION

**State
assumptions,
criteria and the
human owner**

✓ reviewable

EVIDENCE

**Save prompts,
output, checks,
revisions and
approvals**

✓ reviewable

QUALITY

**Pass the
acceptance
criteria and
answer reflection
questions**

✓ reviewable

Stop if the agent or Copilot output cannot be traced, corrected or safely escalated.

Activity 3: Write the Governed HR AI Pilot Charter

Convert an opportunity into a testable pilot charter with go/no-go gates.

CASE

The executive committee will fund only one 90-day HR AI pilot. It expects a clear value case, human controls and an operating owner.

TOOLS

- Word with Microsoft 365 Copilot
- Excel

OUTPUTS

- Pilot charter
- RACI
- Metric baseline and gate table

GUARDRAILS

- Synthetic or de-identified pilot data
- Named risk owner
- Kill switch and rollback

Detailed procedure, prompts and evidence checklist: Learner Guide + activity folder

Activity 3 success map

Show your reasoning—not only a polished output.

INPUT

**Use the supplied
synthetic case
and approved
sources**

✓ reviewable

DECISION

**State
assumptions,
criteria and the
human owner**

✓ reviewable

EVIDENCE

**Save prompts,
output, checks,
revisions and
approvals**

✓ reviewable

QUALITY

**Pass the
acceptance
criteria and
answer reflection
questions**

✓ reviewable

Stop if the agent or Copilot output cannot be traced, corrected or safely escalated.

Unit 1 integration

Connect concept, evidence and action

01 Key decision

Choose the lightest capability that safely solves the work

02 Evidence

What would make the recommendation reviewable?

03 Human role

Who remains accountable?

04 Next move

What will you test or change?

Microsoft 365 Copilot Across the HR Lifecycle

Generate and analyse with a human in control

FRAME

GENERATE


ORCHESTRATE

GOVERN

Concepts → evidence → activity → reflection

A grounded prompt pattern

Context plus task plus constraints plus evidence produces reviewable work.



Design question • What evidence would make this safe enough for HR work?

Workforce planning

Copilot can turn fragmented evidence into an explainable workforce hypothesis.

01

Summarise demand signals

02

Identify capability gaps

03

Draft role and capacity options

04

Keep assumptions visible for challenge

Job design

Better job descriptions start with outcomes and skills, not copied history.

**Translate business need
into outcomes**

**Better job descriptions start
with outcomes and skills, not
copied history.**

**Separate essential from
trainable skills**

**Use inclusive, job-related
language**

**Validate with hiring
manager and policy**

Recruitment campaign

Copilot can adapt one approved proposition across channels without changing facts.

1 Candidate persona and value proposition

2 Channel-specific message


3 Consistent salary and requirements

4 Human approval before publication

Human accountability remains visible at every layer.

Resume screening

AI should organise job-related evidence—not make an opaque hiring decision.



Design question • What evidence would make this safe enough for HR work?

Structured interviewing

A shared evidence model improves consistency more than clever questions alone.

01

Behavioural questions tied to outcomes

02

Anchored scoring guide

03

Independent notes before consensus

04

Candidate accessibility and fairness checks

Candidate communication

Copilot can make communication timely, consistent and humane.

**Personalise only from
approved facts**

**Copilot can make
communication timely,
consistent and humane.**

**Avoid inferred sensitive
attributes**

Give realistic next steps

**Review tone, accuracy
and confidentiality**

Policy drafting


Copilot accelerates comparison and drafting; policy owners approve meaning.

- 1 Ground in current policy sources**
- 2 Surface conflicts and missing definitions**
- 3 Use plain-language examples**
- 4 Preserve version, approver and effective date**

Human accountability remains visible at every layer.

Onboarding

A personalised journey connects role, people, policy, systems and learning.



Design question • What evidence would make this safe enough for HR work?

Performance and development

Copilot can organise evidence and coaching preparation without replacing the conversation.

01

Summarise agreed goals and evidence

02

Separate observation from inference

03

Draft balanced coaching questions

04

Manager validates before discussion

Learning and internal mobility

Skills intelligence becomes useful when employees can challenge and update it.

Map role outcomes to skills

Skills intelligence becomes useful when employees can challenge and update it.

Recommend explainable learning options

Support career conversations

Avoid deterministic potential scores

Engagement and workforce insights

Patterns guide questions; they do not diagnose individuals.

- 1 Aggregate data to protect privacy**
- 2 Triangulate survey, service and retention signals**
- 3 Show uncertainty and limitations**
- 4 Convert insight into a human-led action**

Human accountability remains visible at every layer.

Activity 4: Build a Fair Resume Evidence Matrix

Use Microsoft 365 Copilot to extract job-related evidence and prepare a human-reviewed shortlist.

CASE

HarbourLight received six synthetic applications for the Workforce Operations Analyst role. The hiring manager wants a fast shortlist without an opaque score.

TOOLS

- Excel with Microsoft 365 Copilot
- Word

OUTPUTS

- Evidence matrix
- Uncertainty log
- Human shortlist rationale

GUARDRAILS

- Synthetic resumes only
- Protected attributes excluded
- AI cannot reject candidates

Detailed procedure, prompts and evidence checklist: Learner Guide + activity folder

Activity 4 success map

Show your reasoning—not only a polished output.

INPUT

**Use the supplied
synthetic case
and approved
sources**

✓ reviewable

DECISION

**State
assumptions,
criteria and the
human owner**

✓ reviewable

EVIDENCE

**Save prompts,
output, checks,
revisions and
approvals**

✓ reviewable

QUALITY

**Pass the
acceptance
criteria and
answer reflection
questions**

✓ reviewable

Stop if the agent or Copilot output cannot be traced, corrected or safely escalated.

Activity 5: Draft and Govern an HR Policy Update

Use Word Copilot to compare sources and prepare a controlled policy draft plus communication plan.

CASE

Hybrid-work guidance exists in three conflicting documents. Employees receive inconsistent answers and managers are applying different approval rules.

TOOLS

- Word with Microsoft 365 Copilot
- SharePoint

OUTPUTS

- Conflict register
- Policy draft
- Approval and communication checklist

GUARDRAILS

- Only approved source documents
- Policy owner signs off
- Effective date and superseded versions recorded

Detailed procedure, prompts and evidence checklist: Learner Guide + activity folder

Activity 5 success map

Show your reasoning—not only a polished output.

INPUT

**Use the supplied
synthetic case
and approved
sources**

✓ reviewable

DECISION

**State
assumptions,
criteria and the
human owner**

✓ reviewable

EVIDENCE

**Save prompts,
output, checks,
revisions and
approvals**

✓ reviewable

QUALITY

**Pass the
acceptance
criteria and
answer reflection
questions**

✓ reviewable

Stop if the agent or Copilot output cannot be traced, corrected or safely escalated.

Activity 6: Create a Personalised Onboarding Journey

Use Microsoft 365 Copilot to prepare role-based onboarding plans, communications and learning prompts.

CASE

Three new employees join the distribution hub in different roles and locations. HR needs a consistent core journey with role-specific support.

TOOLS

- Teams
- Word and Outlook with Microsoft 365 Copilot
- Excel

OUTPUTS

- 30-day onboarding plans
- Welcome communications
- Milestone and owner tracker

GUARDRAILS

- Use only approved role and policy data
- Do not infer personal needs
- New joiner can request adjustments

Detailed procedure, prompts and evidence checklist: Learner Guide + activity folder

Activity 6 success map

Show your reasoning—not only a polished output.

INPUT

**Use the supplied
synthetic case
and approved
sources**

✓ reviewable

DECISION

**State
assumptions,
criteria and the
human owner**

✓ reviewable

EVIDENCE

**Save prompts,
output, checks,
revisions and
approvals**

✓ reviewable

QUALITY

**Pass the
acceptance
criteria and
answer reflection
questions**

✓ reviewable

Stop if the agent or Copilot output cannot be traced, corrected or safely escalated.

Unit 2 integration

Connect concept, evidence and action

01 **Key decision**

Generate and analyse with a human in control

02 **Evidence**

What would make the recommendation reviewable?

03 **Human role**

Who remains accountable?

04 **Next move**

What will you test or change?

Build Governed HR Agents

Agent Builder for knowledge; Copilot Studio for orchestration

FRAME

GENERATE


ORCHESTRATE

GOVERN

Concepts → evidence → activity → reflection

Choose the build surface

Agent Builder is fast for team knowledge; Copilot Studio supports broader orchestration.



Design question • What evidence would make this safe enough for HR work?

Agent Builder pattern

A useful team agent has a narrow purpose and authoritative knowledge.

01

Name the audience and task

02

Ground in selected SharePoint sources

03

Add response and refusal instructions

04

Share only with intended users

Copilot Studio pattern

An enterprise agent separates conversation, reasoning and transaction.

**Topic or generative
orchestration**

**An enterprise agent
separates conversation,
reasoning and transaction.**

**Approved knowledge
sources**

Actions and connectors

**Authentication, analytics
and ALM**

Knowledge grounding


The answer should be traceable to the current approved source.

- 1 Curate source ownership**
- 2 Remove obsolete duplicates**
- 3 Respect permissions**
- 4 Return citation, effective date and escalation**

Human accountability remains visible at every layer.

Instructions and boundaries

Agent instructions are an operating policy, not marketing copy.



Design question • What evidence would make this safe enough for HR work?

Tools and agent flows

Tools convert a conversation into work; flows make critical actions deterministic.

01

Use least-privilege connectors

02

Validate inputs before action

03

Add human approval for consequential steps

04

Write outcome and error to an audit log

HR policy agent

A policy agent should answer, cite, qualify and escalate.

Recognise the employee intent

A policy agent should answer, cite, qualify and escalate.

Retrieve the exact current clause

Explain in plain language

Escalate ambiguity or personal exceptions

Onboarding agent


An onboarding agent coordinates milestones without inventing employment commitments.

- 1 Personalise from approved role data**
- 2 Trigger task owners and reminders**
- 3 Track exceptions and access blockers**
- 4 Route sensitive questions to HR**

Human accountability remains visible at every layer.

Leave agent

The agent gathers intent; a governed workflow validates and records the request.



Design question • What evidence would make this safe enough for HR work?

Benefits, payroll and case triage

Self-service is valuable when the agent knows when not to answer.

01

Handle common policy questions

02

Retrieve authorised personal status

03

Never expose another employee's data

04

Escalate disputed or sensitive cases

Offboarding agent

Offboarding requires coordination, access control and respectful communication.

Confirm final date and obligations

Offboarding requires coordination, access control and respectful communication.

Trigger asset and access tasks

Capture knowledge-transfer evidence

Close only after owner confirmations

Agent testing

Test conversations, actions, permissions and recovery—not only happy-path answers.

- 1 Normal and ambiguous intents**
- 2 Prompt injection and data exfiltration**
- 3 Connector failures and duplicate actions**
- 4 Escalation, rollback and audit completeness**

Human accountability remains visible at every layer.

Activity 7: Build an HR Policy Q&A Agent

Create a bounded policy agent in Agent Builder for Microsoft 365 Copilot and test grounded answers.

CASE

Employees ask repetitive hybrid-work and benefits questions in Teams. HR needs a reliable self-service front door that cites policy and escalates exceptions.

TOOLS

- Agent Builder in Microsoft 365 Copilot
- SharePoint
- Teams

OUTPUTS

- Agent purpose and instructions
- Knowledge map
- 12-case test report

GUARDRAILS

- Cite approved policy
- No personalised entitlement decisions
- Escalate ambiguity and sensitive cases

Detailed procedure, prompts and evidence checklist: Learner Guide + activity folder

Activity 7 success map

Show your reasoning—not only a polished output.

INPUT

**Use the supplied
synthetic case
and approved
sources**

✓ reviewable

DECISION

**State
assumptions,
criteria and the
human owner**

✓ reviewable

EVIDENCE

**Save prompts,
output, checks,
revisions and
approvals**

✓ reviewable

QUALITY

**Pass the
acceptance
criteria and
answer reflection
questions**

✓ reviewable

Stop if the agent or Copilot output cannot be traced, corrected or safely escalated.

Activity 8: Design an Onboarding Coordination Agent

Design a Copilot Studio agent that coordinates onboarding while retaining accountable task owners.

CASE

New joiners miss equipment, access and manager check-ins because work is spread across email, service tickets and spreadsheets.

TOOLS

- Microsoft Copilot Studio
- Teams
- SharePoint
- Power Automate or agent flows

OUTPUTS

- Conversation map
- Tool and approval design
- Test evidence

GUARDRAILS

- Least-privilege connectors
- No invented commitments
- Human owner confirms every completion

Detailed procedure, prompts and evidence checklist: Learner Guide + activity folder

Activity 8 success map

Show your reasoning—not only a polished output.

INPUT

**Use the supplied
synthetic case
and approved
sources**

✓ reviewable

DECISION

**State
assumptions,
criteria and the
human owner**

✓ reviewable

EVIDENCE

**Save prompts,
output, checks,
revisions and
approvals**

✓ reviewable

QUALITY

**Pass the
acceptance
criteria and
answer reflection
questions**

✓ reviewable

Stop if the agent or Copilot output cannot be traced, corrected or safely escalated.

Activity 9: Build a Leave Request Agent and Approval Flow

Design a Copilot Studio agent plus deterministic flow for eligibility, approval, update and notification.

CASE

Employees submit leave through messages that HR rekeys into a tracker. Errors occur around balance, overlapping dates and manager approvals.

TOOLS

- Microsoft Copilot Studio
- Agent flows or Power Automate
- Teams
- Excel or HRIS sandbox

OUTPUTS

- Agent instructions
- Flow design
- Approval matrix
- Audit log and test report

GUARDRAILS

- Confirm before submission
- Manager approves exceptions
- Idempotency prevents duplicates

Detailed procedure, prompts and evidence checklist: Learner Guide + activity folder

Activity 9 success map

Show your reasoning—not only a polished output.

INPUT

**Use the supplied
synthetic case
and approved
sources**

✓ reviewable

DECISION

**State
assumptions,
criteria and the
human owner**

✓ reviewable

EVIDENCE

**Save prompts,
output, checks,
revisions and
approvals**

✓ reviewable

QUALITY

**Pass the
acceptance
criteria and
answer reflection
questions**

✓ reviewable

Stop if the agent or Copilot output cannot be traced, corrected or safely escalated.

Unit 3 integration

Connect concept, evidence and action

01 **Key decision**

Agent Builder for knowledge; Copilot Studio for orchestration

02 **Evidence**

What would make the recommendation reviewable?

03 **Human role**

Who remains accountable?

04 **Next move**

What will you test or change?

LEARNING UNIT 4

Responsible Adoption and HR Transformation

Govern decisions, data, people and agent operations

FRAME

GENERATE


ORCHESTRATE

GOVERN

Concepts → evidence → activity → reflection

Why HR AI is high consequence

Employment decisions affect opportunity, income, dignity and trust.



Design question • What evidence would make this safe enough for HR work?

Human involvement by impact

More consequential and less reversible decisions need stronger human authority.

01

Inform: AI drafts or retrieves

02

Recommend: AI proposes with evidence

03

Approve: authorised human decides

04

Execute: agent acts only within approved bounds

Privacy by design

Minimise data before asking a model to protect it.

Purpose and lawful use

**Minimise data before asking a
model to protect it.**

**Minimum necessary
fields**

**Access, retention and
deletion**

**Incident and data-subject
response**

Fair hiring controls

Selection criteria must remain job-related, reviewable and consistently applied.

- 1 Fairly consider candidates**
- 2 Exclude non-job-related protected traits**
- 3 Keep interview and decision evidence**
- 4 Provide human challenge and correction**

Human accountability remains visible at every layer.

Bias is a system property

Bias can enter through objectives, data, prompts, tools and human use.



Design question • What evidence would make this safe enough for HR work?

Transparency and explainability

Employees need to know when AI is used and how to seek help.

01

Disclose the AI role

02

Show source and material factors

03

Explain limits and uncertainty

04

Provide an accessible human channel

Security and prompt injection

An agent can be manipulated through the text it reads.

**Treat retrieved content
as untrusted**

**An agent can be manipulated
through the text it reads.**

**Separate instructions
from data**

**Allow-list tools and
destinations**

**Monitor anomalous
actions**

Agent identity and least privilege


Every action should be attributable to an authorised identity.

- 1 Use named service or user context**
- 2 Grant only required permissions**
- 3 Separate maker, approver and operator roles**
- 4 Review access and dormant agents**

Human accountability remains visible at every layer.

Operational controls

Agents are digital workers that need owners, logs, limits and retirement.



Design question • What evidence would make this safe enough for HR work?

Success measures

Measure service, adoption, quality, risk and human value together.

01

Cycle time and first-contact resolution

02

Accuracy and evidence completeness

03

Escalation and override rates

04

Employee trust and time returned to people

Change and AI literacy

Adoption depends on role clarity and practical learning.

**Explain what changes
and what does not**

**Adoption depends on role
clarity and practical learning.**

**Train prompting, review
and escalation**

**Create champions and
feedback loops**

**Recognise new work:
agent owner, tester,
steward**

90-day responsible roadmap

Move from a bounded pilot to an owned capability through evidence gates.

- 1 Days 1–30: baseline, risk and prototype**
- 2 Days 31–60: test, controls and user feedback**
- 3 Days 61–90: measured pilot and operating handover**
- 4 Scale only when value and control thresholds pass**

Human accountability remains visible at every layer.

Activity 10: Prepare a Fair Performance and Learning Conversation

Use Copilot to organise evidence, distinguish observation from inference and prepare human-led coaching.

CASE

A manager has fragmented notes, goals and service metrics for two employees and wants Copilot to draft development conversations.

TOOLS

- Word and Teams with Microsoft 365 Copilot
- Excel

OUTPUTS

- Evidence summary
- Coaching questions
- Learning plan

GUARDRAILS

- No personality inference
- Employee can correct evidence
- Manager owns the conversation and rating

Detailed procedure, prompts and evidence checklist: Learner Guide + activity folder

Activity 10 success map

Show your reasoning—not only a polished output.

INPUT

**Use the supplied
synthetic case
and approved
sources**

✓ reviewable

DECISION

**State
assumptions,
criteria and the
human owner**

✓ reviewable

EVIDENCE

**Save prompts,
output, checks,
revisions and
approvals**

✓ reviewable

QUALITY

**Pass the
acceptance
criteria and
answer reflection
questions**

✓ reviewable

Stop if the agent or Copilot output cannot be traced, corrected or safely escalated.

Activity 11: Design Safe Employee Service and Offboarding

Define privacy-aware routing, escalation and least-privilege actions across sensitive HR services.

CASE

HR wants one conversational front door for benefits, payroll status, employee-relations triage and offboarding tasks.

TOOLS

- Microsoft Copilot Studio
- Teams
- SharePoint
- HRIS sandbox

OUTPUTS

- Intent-and-risk catalogue
- Escalation matrix
- Offboarding control map

GUARDRAILS

- Authenticate personal requests
- No legal or clinical advice
- Separation of duties for access removal

Detailed procedure, prompts and evidence checklist: Learner Guide + activity folder

Activity 11 success map

Show your reasoning—not only a polished output.

INPUT

**Use the supplied
synthetic case
and approved
sources**

✓ reviewable

DECISION

**State
assumptions,
criteria and the
human owner**

✓ reviewable

EVIDENCE

**Save prompts,
output, checks,
revisions and
approvals**

✓ reviewable

QUALITY

**Pass the
acceptance
criteria and
answer reflection
questions**

✓ reviewable

Stop if the agent or Copilot output cannot be traced, corrected or safely escalated.

Activity 12: Capstone: Operate the HR Agent Control Room

Use service, risk, adoption and human-value evidence to make a defensible operating decision.

CASE

HarbourLight's pilot portfolio now includes policy, onboarding and leave agents. Leaders must decide whether to scale, remediate or stop each agent.

TOOLS

- Excel
- Teams
- Microsoft 365 Copilot

OUTPUTS

- Agent scorecard
- Incident response
- 90-day roadmap and executive recommendation

GUARDRAILS

- No scale without control thresholds
- Employee feedback included
- Named owner and retirement path

Detailed procedure, prompts and evidence checklist: Learner Guide + activity folder

Activity 12 success map

Show your reasoning—not only a polished output.

INPUT

**Use the supplied
synthetic case
and approved
sources**

✓ reviewable

DECISION

**State
assumptions,
criteria and the
human owner**

✓ reviewable

EVIDENCE

**Save prompts,
output, checks,
revisions and
approvals**

✓ reviewable

QUALITY

**Pass the
acceptance
criteria and
answer reflection
questions**

✓ reviewable

Stop if the agent or Copilot output cannot be traced, corrected or safely escalated.

Unit 4 integration

Connect concept, evidence and action

01

Key decision

Govern decisions, data, people and agent operations

02

Evidence

What would make the recommendation reviewable?

03

Human role

Who remains accountable?


04

Next move

What will you test or change?

A balanced HR AI scorecard

Speed is useful only when quality, trust and control remain healthy.



READ TOGETHER

A lower cycle time cannot compensate for weak evidence, falling trust or missing approvals. Pair one value metric with at least one quality and one risk guardrail.

HR lifecycle coverage

One operating system, many employee moments

01 Attract

Workforce planning, job design, recruitment and selection

02 Join

Onboarding, access, policy and manager connection

03 Grow

Performance, learning, mobility and engagement

04 Serve

Leave, benefits, payroll and employee relations

05 Exit

Offboarding, access, assets and knowledge transfer

06 Operate

Metrics, incidents, governance and retirement

The responsible-agent operating loop

Agents remain a managed organisational capability

01 1. Observe

Service, quality, adoption and risk signals

02 2. Review

Human owner inspects evidence and exceptions

03 3. Improve

Update knowledge, instructions, tools or process

04 4. Validate

Regression, fairness, security and user testing

05 5. Decide

Scale, remediate, pause or retire

Your 90-day pathway

Move through evidence gates

01 Days 1–30

Baseline • use-case boundary •
privacy/risk review • prototype

02 Days 31–60

Test • human controls • user
feedback • operating design

03 Days 61–90

Measured pilot • incident
rehearsal • scale decision •
handover

Key takeaways

Create with Copilot • orchestrate with agents • govern as HR

01 Choose

Use the lightest sufficient capability.

02 Ground

Use authoritative knowledge and job-related evidence.

03 Control

Keep human authority for consequential decisions.

04 Operate

Measure value, quality, trust and risk together.

Assessment reminder

Demonstrate your own competence

01 Written Assessment

60 minutes • K1–K5

02 Case Study Assessment

60 minutes • A1–A5

03 Before submission

Answer every part • check evidence • protect confidentiality

Assessment flow

Read → plan → answer → verify → submit

01 Read

Underline the command word.

02 Plan

Use time deliberately.

03 Answer

Apply course concepts.

04 Verify

Check completeness and clarity.

05 Submit

Follow assessor directions.

Digital attendance

Complete the final attendance check when directed

01 Check

Use the correct session and learner identity.

02 Confirm

Wait for the successful confirmation.

03 Support

Tell the trainer immediately if the check fails.

Thank you

Design HR AI that gives people better service and humans better judgement

01 Course

Agentic AI for HR

02 Code

TGS-2024045795

03 Resources

<https://github.com/tertiarycourses/TGS-2024045795-Agentic-AI-for-HR>