

TERTIARY INFOTECH ACADEMY

LEARNER GUIDE

Agentic AI for HR

Course Code: TGS-2024045795

2 days / 16 hours

Human Resource Digitalisation-4 (HRS-HRM-4031-1.1)

Version 6.0 • 16 August 2026

How to use this guide

The slide deck explains concepts visually. This Learner Guide contains the detailed activity procedures, prompts, evidence requirements and questions. All cases and records are synthetic for learning; do not substitute live employee data.

Course information

Field	Detail
Course	Agentic AI for HR
Course code	TGS-2024045795
Duration	2 days / 16 hours
TSC	Human Resource Digitalisation-4 (HRS-HRM-4031-1.1)
Trainer	Dr. Alfred Ang

Learning outcomes

- LO1. Compare how generative and agentic AI reshape HR operating models versus conventional HR delivery.
- LO2. Analyse where Microsoft 365 Copilot and AI agents can improve end-to-end HR processes and employee experience.
- LO3. Conduct feasibility analysis and design a governed implementation for an HR generative-AI or agentic-AI use case.
- LO4. Apply legal, ethical, privacy, security, fairness and human-oversight controls to AI-enabled HR work.

Platform model

Need	Preferred Microsoft capability	Human role
Draft, summarise, compare or analyse	Microsoft 365 Copilot in Word, Excel, Outlook, Teams	Provide context, verify evidence and approve output
Focused team Q&A over approved knowledge	Agent Builder in Microsoft 365 Copilot	Curate knowledge, test and manage sharing
Multi-step workflow across systems	Microsoft Copilot Studio with tools/agent flows	Approve consequential actions and operate controls

Contents and slide alignment

Unit	Coverage	Slides
1	From Generative AI to Agentic HR	17–36
2	Microsoft 365 Copilot Across the HR Lifecycle	37–56
3	Build Governed HR Agents	57–76
4	Responsible Adoption and HR	77–96

	Transformation	
--	----------------	--

Learning Unit 1: From Generative AI to Agentic HR

Choose the lightest capability that safely solves the work

Concept notes

The HR AI continuum

Automation follows rules; copilots create; agents pursue bounded goals.

- Rules automate predictable hand-offs
- Microsoft 365 Copilot drafts, summarises and analyses
- Agents reason, plan, use tools and act
- Humans retain accountability

Sources: <https://www.ibm.com/think/topics/ai-agents-in-human-resources>; <https://www.hrci.org/community/blogs-and-announcements/hr-leads-business-blog/hr-leads-business/2025/08/21/agentic-ai-in-hr>

Conventional HR operating model

Queues and hand-offs make service slow and opaque.

- Requests arrive through email and forms
- Specialists search systems manually
- Approvals cross organisational boundaries
- Evidence is scattered across tools

Sources: <https://www.ibm.com/think/topics/ai-in-hr>; <https://www.zendesk.com/blog/ai/workflow-automation/ai-in-hr/>

Copilot-assisted HR operating model

A human remains the operator while Copilot accelerates knowledge work.

- Ground prompts in approved work content
- Generate a first draft, not a final decision
- Compare and refine with domain expertise
- Record source and reviewer

Sources: <https://www.microsoft.com/en-us/microsoft-copilot/copilot-101/ai-for-hr>;
<https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/the-organization-blog/four-ways-to-start-using-generative-ai-in-hr>

Agent-enabled HR operating model

An agent can coordinate a bounded workflow across systems.

- Interpret an employee intent
- Retrieve policy and contextual data
- Call approved tools and workflows
- Escalate exceptions with an audit trail

Sources: <https://www.ibm.com/think/topics/ai-agents-in-human-resources>;
<https://www.moveworks.com/us/en/resources/blog/agentic-ai-in-hr>

Assistant, agent or workflow?

Use the simplest lever that meets the outcome and risk profile.

- Copilot: create or analyse with a person
- Agent: converse, reason and orchestrate
- Flow: deterministic transaction

- Hybrid: agent front door plus governed flow

Sources: <https://www.strategyand.pwc.com/de/en/functions/organisational-strategy/agentic-ai-in-hr.html>;
<https://www.microsoft.com/en-sg/microsoft-365-copilot/microsoft-copilot-studio>

The Microsoft HR AI stack

Microsoft 365 Copilot and Copilot Studio play different but complementary roles.

- Word and Outlook: content and communication
- Excel: evidence and analysis
- Teams and SharePoint: collaboration and knowledge
- Copilot Studio: agents, tools, flows and channels

Sources: <https://www.microsoft.com/en-us/microsoft-copilot/copilot-101/ai-for-hr>; <https://www.microsoft.com/en-sg/microsoft-365-copilot/microsoft-copilot-studio>

Agent anatomy

Reliable agents need a goal, knowledge, instructions, tools, boundaries and tests.

- Goal defines the employee outcome
- Knowledge constrains grounded answers
- Instructions set behaviour and refusal rules
- Tools act; tests prove expected behaviour

Sources: <https://www.ibm.com/think/topics/ai-agents-in-human-resources>;
<https://www.awinfracloud.com/blog/build-hr-agent-copilot-studio>

HR lifecycle opportunity map

AI value appears across attract, join, grow, serve and exit.

- Attract: workforce planning and recruitment
- Join: onboarding and access
- Grow: performance, learning and mobility
- Serve and exit: policy, leave, cases and offboarding

Sources: <https://www.sap.com/resources/ai-for-hr>; <https://www.ibm.com/think/topics/ai-agents-in-human-resources>

Value without losing the human

Saved time should be redirected to judgement, coaching and care.

- Reduce repetitive search and drafting
- Improve access and consistency
- Create capacity for human conversations
- Do not automate empathy or accountability

Sources: <https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/the-organization-blog/four-ways-to-start-using-generative-ai-in-hr>; <https://www.auxis.com/ai-in-hr-examples-types-use-cases-best-practices/>

Use-case selection

Start with frequent, bounded and measurable work.

- High volume and clear policy
- Reliable knowledge and integrations
- Reversible actions and escalation
- A KPI baseline and accountable owner

Sources: <https://www.gartner.com/en/human-resources/topics/artificial-intelligence-in-hr>;
<https://www.auxis.com/ai-in-hr-examples-types-use-cases-best-practices/>

Feasibility lens

A good idea is not yet a deployable use case.

- Business value and employee impact
- Data readiness and permissions
- Technical reach and failure modes
- Risk, change and operating ownership

Sources: <https://www.gartner.com/en/human-resources/topics/artificial-intelligence-in-hr>;
<https://www.strategyand.pwc.com/de/en/functions/organisational-strategy/agentic-ai-in-hr.html>

From pilot to operating capability

A pilot must prove usefulness, control and adoption together.

- Define baseline and success threshold
- Use synthetic data before live data
- Test normal, edge and adversarial cases
- Assign monitoring, review and retirement

Sources: <https://www.auxis.com/ai-in-hr-examples-types-use-cases-best-practices/>;
<https://www.strategyand.pwc.com/de/en/functions/organisational-strategy/agentic-ai-in-hr.html>

Activity 1: Map the HR AI Opportunity Portfolio

Scenario

HarbourLight Logistics SG has 620 employees, a five-person HR team and rising service backlogs. The CHRO wants a defensible starting portfolio—not a list of fashionable tools.

Learning goal

Classify HR work across automation, Microsoft 365 Copilot and agentic AI, then select one bounded pilot.

Tools

Microsoft Teams; Microsoft 365 Copilot; Excel

Procedure

1. Review the synthetic HR demand and pain-point data in the workbook.
2. Define the employee or manager outcome for each use case.
3. Classify the work as deterministic automation, Copilot-assisted knowledge work, agentic orchestration or unsuitable.
4. Score frequency, value, data readiness, reversibility and people risk using the workbook rubric.
5. Use Copilot to challenge the top three options and list hidden dependencies.
6. Select one pilot and one explicit non-goal.
7. Record baseline, success threshold, owner and stop condition.
8. Peer-review another team's selection for over-automation or missing stakeholders.

Prompt scaffold

You are supporting the Agentic AI for HR case. Use only the supplied synthetic evidence. Task: Classify HR work across automation, Microsoft 365 Copilot and agentic AI, then select one bounded pilot. State assumptions, cite the evidence used, distinguish observation from inference, identify uncertainty, and list the human checks required before use.

Evidence to submit

- Completed opportunity-and-risk matrix
- Two-minute pilot recommendation
- Evidence log

Acceptance criteria

- No employee personal data
- State all assumptions
- Human owner for every shortlisted use case

The submission is complete when the workbook evidence log contains the input/source, prompt or design decision, output, verification, revision and named human owner.

Scenario questions

5. Why is the selected capability the lightest sufficient option?
6. What evidence would stop the pilot?
7. Whose experience could worsen even if cycle time improves?

Activity 2: Design a Skills-Based Workforce Plan

Scenario

HarbourLight is opening a digitally enabled distribution hub. HR must decide whether to hire, redeploy or develop people for a new Workforce Operations Analyst role.

Learning goal

Use Copilot to synthesise demand evidence and draft a skills-based, fair job design.

Tools

Excel with Microsoft 365 Copilot; Word with Microsoft 365 Copilot

Procedure

1. Inspect the synthetic demand, workload and skills inventory tabs.
2. Ask Copilot to summarise demand drivers and cite the cells used.
3. Draft hire, redeploy and develop options with trade-offs.
4. Choose a planning hypothesis and mark uncertainty.
5. Translate role outcomes into essential and trainable skills.
6. Draft a job description in Word grounded in the approved skills table.
7. Run an inclusive-language and job-relatedness review.
8. Document what the hiring manager must validate.

Prompt scaffold

You are supporting the Agentic AI for HR case. Use only the supplied synthetic evidence. Task: Use Copilot to synthesise demand evidence and draft a skills-based, fair job design. State assumptions, cite the evidence used, distinguish observation from inference, identify uncertainty, and list the human checks required before use.

Evidence to submit

- Workforce hypothesis
- Skills gap table
- Inclusive job description

Acceptance criteria

- Separate evidence from assumption
- No age, nationality or other non-job-related criteria
- Hiring manager approval

The submission is complete when the workbook evidence log contains the input/source, prompt or design decision, output, verification, revision and named human owner.

Scenario questions

8. Which assumption has the greatest planning impact?
9. Which requirement could unfairly narrow the pool?
10. How will the hypothesis be revisited?

Activity 3: Write the Governed HR AI Pilot Charter

Scenario

The executive committee will fund only one 90-day HR AI pilot. It expects a clear value case, human controls and an operating owner.

Learning goal

Convert an opportunity into a testable pilot charter with go/no-go gates.

Tools

Word with Microsoft 365 Copilot; Excel

Procedure

1. Select the use case approved in Activity 1 or use the provided leave-agent case.
2. Define the user, outcome, trigger, boundary and non-goals.
3. Map data, knowledge, tools, decisions and human approvals.
4. Ask Copilot to draft a charter from the workbook evidence.
5. Add baseline, target, guardrail and stop metrics.
6. Assign business owner, technical owner, policy owner, approver and support roles.
7. Specify pilot population, test data and feedback route.
8. Run a red-team review and revise the charter.

Prompt scaffold

You are supporting the Agentic AI for HR case. Use only the supplied synthetic evidence. Task: Convert an opportunity into a testable pilot charter with go/no-go gates. State assumptions, cite the evidence used, distinguish observation from inference, identify uncertainty, and list the human checks required before use.

Evidence to submit

- Pilot charter
- RACI
- Metric baseline and gate table

Acceptance criteria

- Synthetic or de-identified pilot data
- Named risk owner
- Kill switch and rollback

The submission is complete when the workbook evidence log contains the input/source, prompt or design decision, output, verification, revision and named human owner.

Scenario questions

11. Which decision remains human and why?
12. Who can stop the pilot?
13. What would a misleading success metric reward?

Learning Unit 2: Microsoft 365 Copilot Across the HR Lifecycle

Generate and analyse with a human in control

Concept notes

A grounded prompt pattern

Context plus task plus constraints plus evidence produces reviewable work.

- Role and business context
- Specific task and desired outcome
- Approved sources and boundaries
- Output format and verification checks

Sources: <https://www.microsoft.com/en-us/microsoft-copilot/copilot-101/ai-for-hr>

Workforce planning

Copilot can turn fragmented evidence into an explainable workforce hypothesis.

- Summarise demand signals
- Identify capability gaps
- Draft role and capacity options
- Keep assumptions visible for challenge

Sources: <https://www.gartner.com/en/human-resources/topics/artificial-intelligence-in-hr>;
<https://www.microsoft.com/en-us/microsoft-copilot/copilot-101/ai-for-hr>

Job design

Better job descriptions start with outcomes and skills, not copied history.

- Translate business need into outcomes
- Separate essential from trainable skills
- Use inclusive, job-related language
- Validate with hiring manager and policy

Sources: <https://www.microsoft.com/en-us/microsoft-copilot/copilot-101/ai-for-hr>;
<https://www.mom.gov.sg/employment-practices/fair-consideration-framework>

Recruitment campaign

Copilot can adapt one approved proposition across channels without changing facts.

- Candidate persona and value proposition
- Channel-specific message
- Consistent salary and requirements
- Human approval before publication

Sources: <https://www.microsoft.com/en-us/microsoft-copilot/copilot-101/ai-for-hr>;
<https://www.sap.com/resources/ai-for-hr>

Resume screening

AI should organise job-related evidence—not make an opaque hiring decision.

- Build a disclosed scoring rubric
- Extract evidence, uncertainty and gaps

- Blind protected attributes where feasible
- A human reviews every shortlist decision

Sources: <https://www.microsoft.com/en-us/microsoft-copilot/copilot-101/ai-for-hr>;
<https://www.mom.gov.sg/employment-practices/fair-consideration-framework>

Structured interviewing

A shared evidence model improves consistency more than clever questions alone.

- Behavioural questions tied to outcomes
- Anchored scoring guide
- Independent notes before consensus
- Candidate accessibility and fairness checks

Sources: <https://www.sap.com/resources/ai-for-hr>; <https://www.pdpc.gov.sg/help-and-resources/2020/01/model-ai-governance-framework>

Candidate communication

Copilot can make communication timely, consistent and humane.

- Personalise only from approved facts
- Avoid inferred sensitive attributes
- Give realistic next steps
- Review tone, accuracy and confidentiality

Sources: <https://www.microsoft.com/en-us/microsoft-copilot/copilot-101/ai-for-hr>; <https://www.pdpc.gov.sg/help-and-resources/2020/01/model-ai-governance-framework>

Policy drafting

Copilot accelerates comparison and drafting; policy owners approve meaning.

- Ground in current policy sources
- Surface conflicts and missing definitions
- Use plain-language examples
- Preserve version, approver and effective date

Sources: <https://www.microsoft.com/en-us/microsoft-copilot/copilot-101/ai-for-hr>;
<https://www.zendesk.com/blog/ai/workflow-automation/ai-in-hr/>

Onboarding

A personalised journey connects role, people, policy, systems and learning.

- Prepare a role-specific plan
- Summarise must-know policies
- Identify collaboration network
- Track access, milestones and feedback

Sources: <https://adoption.microsoft.com/en-us/scenario-library/human-resources/improve-onboarding-and-development-processes-copilot-studio/>; <https://www.sap.com/resources/ai-for-hr>

Performance and development

Copilot can organise evidence and coaching preparation without replacing the conversation.

- Summarise agreed goals and evidence
- Separate observation from inference
- Draft balanced coaching questions
- Manager validates before discussion

Sources: <https://www.microsoft.com/en-us/microsoft-copilot/copilot-101/ai-for-hr>;
<https://www.sap.com/resources/ai-for-hr>

Learning and internal mobility

Skills intelligence becomes useful when employees can challenge and update it.

- Map role outcomes to skills
- Recommend explainable learning options
- Support career conversations
- Avoid deterministic potential scores

Sources: <https://www.gartner.com/en/human-resources/topics/artificial-intelligence-in-hr>;
<https://www.sap.com/resources/ai-for-hr>

Engagement and workforce insights

Patterns guide questions; they do not diagnose individuals.

- Aggregate data to protect privacy
- Triangulate survey, service and retention signals
- Show uncertainty and limitations
- Convert insight into a human-led action

Sources: <https://www.gartner.com/en/human-resources/topics/artificial-intelligence-in-hr>;
<https://www.pdpc.gov.sg/help-and-resources/2020/01/model-ai-governance-framework>

Activity 4: Build a Fair Resume Evidence Matrix

Scenario

HarbourLight received six synthetic applications for the Workforce Operations Analyst role. The hiring manager wants a fast shortlist without an opaque score.

Learning goal

Use Microsoft 365 Copilot to extract job-related evidence and prepare a human-reviewed shortlist.

Tools

Excel with Microsoft 365 Copilot; Word

Procedure

1. Read the approved job outcomes and anchored rubric.
2. Review the six synthetic resumes without adding criteria.
3. Prompt Copilot to extract verbatim evidence for each criterion and mark gaps as unknown.
4. Check every extracted item against the source resume.
5. Calculate rubric totals using workbook formulas; do not let Copilot invent a score.
6. Review for proxy variables and inconsistent evidence standards.
7. Record human shortlist decisions and reasons separately from the AI extraction.
8. Draft candidate next-step messages without disclosing other candidates' information.

Prompt scaffold

You are supporting the Agentic AI for HR case. Use only the supplied synthetic evidence. Task: Use Microsoft 365 Copilot to extract job-related evidence and prepare a human-reviewed shortlist. State assumptions, cite the evidence used, distinguish observation from inference, identify uncertainty, and list the human checks required before use.

Evidence to submit

- Evidence matrix
- Uncertainty log
- Human shortlist rationale

Acceptance criteria

- Synthetic resumes only
- Protected attributes excluded
- AI cannot reject candidates

The submission is complete when the workbook evidence log contains the input/source, prompt or design decision, output, verification, revision and named human owner.

Scenario questions

14. Where did missing evidence risk becoming a negative assumption?
15. What would make this screening process contestable?
16. Which errors matter most: false inclusion or false exclusion?

Activity 5: Draft and Govern an HR Policy Update

Scenario

Hybrid-work guidance exists in three conflicting documents. Employees receive inconsistent answers and managers are applying different approval rules.

Learning goal

Use Word Copilot to compare sources and prepare a controlled policy draft plus communication plan.

Tools

Word with Microsoft 365 Copilot; SharePoint

Procedure

1. Review the three synthetic policy extracts and source metadata.
2. Ask Copilot to identify conflicts, gaps and undefined terms with citations.
3. Confirm the policy hierarchy with the scenario brief.
4. Draft a plain-language policy using the approved rule set.
5. Add examples, exceptions and an escalation route.
6. Run a privacy, fairness and operational feasibility review.
7. Create the approval, publication and acknowledgement plan.
8. Archive superseded content and record the version change.

Prompt scaffold

You are supporting the Agentic AI for HR case. Use only the supplied synthetic evidence. Task: Use Word Copilot to compare sources and prepare a controlled policy draft plus communication plan. State assumptions, cite the evidence used, distinguish observation from inference, identify uncertainty, and list the human checks required before use.

Evidence to submit

- Conflict register
- Policy draft
- Approval and communication checklist

Acceptance criteria

- Only approved source documents
- Policy owner signs off
- Effective date and superseded versions recorded

The submission is complete when the workbook evidence log contains the input/source, prompt or design decision, output, verification, revision and named human owner.

Scenario questions

17. Which conflict could create unequal treatment?
18. What must not be inferred by Copilot?
19. How will users know which version is authoritative?

Activity 6: Create a Personalised Onboarding Journey

Scenario

Three new employees join the distribution hub in different roles and locations. HR needs a consistent core journey with role-specific support.

Learning goal

Use Microsoft 365 Copilot to prepare role-based onboarding plans, communications and learning prompts.

Tools

Teams; Word and Outlook with Microsoft 365 Copilot; Excel

Procedure

1. Review the three synthetic personas, role outcomes and mandatory controls.
2. Define a common Day 1, Week 1 and Days 8–30 backbone.
3. Prompt Copilot to personalise tasks from role, location and approved learning needs.
4. Check access, safety, policy, network and learning coverage.
5. Draft manager and new-joiner communications in Outlook.
6. Assign owners and due dates in the workbook.
7. Add a feedback pulse and exception route.
8. Compare journeys for consistency and unintended disadvantage.

Prompt scaffold

You are supporting the Agentic AI for HR case. Use only the supplied synthetic evidence. Task: Use Microsoft 365 Copilot to prepare role-based onboarding plans, communications and learning prompts. State assumptions, cite the evidence used, distinguish observation from inference, identify uncertainty, and list the human checks required before use.

Evidence to submit

- 30-day onboarding plans
- Welcome communications
- Milestone and owner tracker

Acceptance criteria

- Use only approved role and policy data
- Do not infer personal needs
- New joiner can request adjustments

The submission is complete when the workbook evidence log contains the input/source, prompt or design decision, output, verification, revision and named human owner.

Scenario questions

20. What should remain standard for every new joiner?
21. Which personalisation requires consent?
22. How will HR detect an onboarding access failure early?

Learning Unit 3: Build Governed HR Agents

Agent Builder for knowledge; Copilot Studio for orchestration

Concept notes

Choose the build surface

Agent Builder is fast for team knowledge; Copilot Studio supports broader orchestration.

- Agent Builder: focused Q&A and team productivity
- Copilot Studio: tools, flows, channels and lifecycle
- Copy an agent to extend it
- Licensing and governance remain design constraints

Sources: <https://www.microsoft.com/en-sg/microsoft-365-copilot/microsoft-copilot-studio>;
<https://www.microsoft.com/en-us/microsoft-copilot/blog/copilot-studio/transform-hr-with-ai-powered-agents/>

Agent Builder pattern

A useful team agent has a narrow purpose and authoritative knowledge.

- Name the audience and task
- Ground in selected SharePoint sources
- Add response and refusal instructions
- Share only with intended users

Sources: <https://adoption.microsoft.com/en-us/scenario-library/human-resources/improve-onboarding-and-development-processes-copilot-studio/>; <https://www.microsoft.com/en-sg/microsoft-365-copilot/microsoft-copilot-studio>

Copilot Studio pattern

An enterprise agent separates conversation, reasoning and transaction.

- Topic or generative orchestration
- Approved knowledge sources
- Actions and connectors
- Authentication, analytics and ALM

Sources: <https://www.microsoft.com/en-sg/microsoft-365-copilot/microsoft-copilot-studio>;
<https://www.awinfracloud.com/blog/build-hr-agent-copilot-studio>

Knowledge grounding

The answer should be traceable to the current approved source.

- Curate source ownership
- Remove obsolete duplicates
- Respect permissions
- Return citation, effective date and escalation

Sources: <https://www.awinfracloud.com/blog/build-hr-agent-copilot-studio>; <https://adoption.microsoft.com/en-us/ai-agents/employee-self-service-agent/>

Instructions and boundaries

Agent instructions are an operating policy, not marketing copy.

- Define in-scope intents
- State prohibited decisions

- Specify clarification and escalation
- Require confirmation before consequential actions

Sources: <https://www.ibm.com/think/topics/ai-agents-in-human-resources>; <https://www.pdpc.gov.sg/help-and-resources/2020/01/model-ai-governance-framework>

Tools and agent flows

Tools convert a conversation into work; flows make critical actions deterministic.

- Use least-privilege connectors
- Validate inputs before action
- Add human approval for consequential steps
- Write outcome and error to an audit log

Sources: <https://www.microsoft.com/en-us/microsoft-copilot/blog/copilot-studio/transform-hr-with-ai-powered-agents/>; <https://adoption.microsoft.com/en-us/ai-agents/employee-self-service-agent/>

HR policy agent

A policy agent should answer, cite, qualify and escalate.

- Recognise the employee intent
- Retrieve the exact current clause
- Explain in plain language
- Escalate ambiguity or personal exceptions

Sources: <https://www.awinfracloud.com/blog/build-hr-agent-copilot-studio/>; <https://www.zendesk.com/blog/ai/workflow-automation/ai-in-hr/>

Onboarding agent

An onboarding agent coordinates milestones without inventing employment commitments.

- Personalise from approved role data
- Trigger task owners and reminders
- Track exceptions and access blockers
- Route sensitive questions to HR

Sources: <https://adoption.microsoft.com/en-us/scenario-library/human-resources/improve-onboarding-and-development-processes-copilot-studio/>; <https://www.microsoft.com/en-us/microsoft-copilot/blog/copilot-studio/transform-hr-with-ai-powered-agents/>

Leave agent

The agent gathers intent; a governed workflow validates and records the request.

- Check balance and policy eligibility
- Confirm dates and coverage
- Route manager approval
- Update system and notify with audit evidence

Sources: <https://adoption.microsoft.com/en-us/ai-agents/employee-self-service-agent/>; <https://www.microsoft.com/en-us/microsoft-copilot/blog/copilot-studio/transform-hr-with-ai-powered-agents/>

Benefits, payroll and case triage

Self-service is valuable when the agent knows when not to answer.

- Handle common policy questions
- Retrieve authorised personal status
- Never expose another employee's data

- Escalate disputed or sensitive cases

Sources: <https://www.zendesk.com/blog/ai/workflow-automation/ai-in-hr/>; <https://www.ibm.com/think/topics/ai-agents-in-human-resources>

Offboarding agent

Offboarding requires coordination, access control and respectful communication.

- Confirm final date and obligations
- Trigger asset and access tasks
- Capture knowledge-transfer evidence
- Close only after owner confirmations

Sources: <https://www.ibm.com/think/topics/ai-agents-in-human-resources>;
<https://www.moveworks.com/us/en/resources/blog/agentic-ai-in-hr>

Agent testing

Test conversations, actions, permissions and recovery—not only happy-path answers.

- Normal and ambiguous intents
- Prompt injection and data exfiltration
- Connector failures and duplicate actions
- Escalation, rollback and audit completeness

Sources: <https://www.pdpc.gov.sg/help-and-resources/2020/01/model-ai-governance-framework>;
<https://www.auxis.com/ai-in-hr-examples-types-use-cases-best-practices/>

Activity 7: Build an HR Policy Q&A Agent

Scenario

Employees ask repetitive hybrid-work and benefits questions in Teams. HR needs a reliable self-service front door that cites policy and escalates exceptions.

Learning goal

Create a bounded policy agent in Agent Builder for Microsoft 365 Copilot and test grounded answers.

Tools

Agent Builder in Microsoft 365 Copilot; SharePoint; Teams

Procedure

1. Define the audience, in-scope intents and explicit exclusions.
2. Select the approved synthetic policy knowledge files.
3. Write instructions for citation, effective date, clarification, refusal and escalation.
4. Create the agent in Agent Builder and add the knowledge sources.
5. Add conversation starters for common employee intents.
6. Run normal, ambiguous, outdated-policy and sensitive-personal-case tests.
7. Record answer, citation, pass/fail and corrective action.
8. Share only with the test group after the policy owner approves.

Prompt scaffold

You are supporting the Agentic AI for HR case. Use only the supplied synthetic evidence. Task: Create a bounded policy agent in Agent Builder for Microsoft 365 Copilot and test grounded answers. State assumptions, cite the evidence used, distinguish observation from inference, identify uncertainty, and list the human checks required before use.

Evidence to submit

- Agent purpose and instructions
- Knowledge map
- 12-case test report

Acceptance criteria

- Cite approved policy
- No personalised entitlement decisions
- Escalate ambiguity and sensitive cases

The submission is complete when the workbook evidence log contains the input/source, prompt or design decision, output, verification, revision and named human owner.

Scenario questions

23. What makes an answer grounded rather than merely plausible?
24. When should the agent stop answering?
25. What new operating work does this agent create?

Activity 8: Design an Onboarding Coordination Agent

Scenario

New joiners miss equipment, access and manager check-ins because work is spread across email, service tickets and spreadsheets.

Learning goal

Design a Copilot Studio agent that coordinates onboarding while retaining accountable task owners.

Tools

Microsoft Copilot Studio; Teams; SharePoint; Power Automate or agent flows

Procedure

1. Map the onboarding trigger, milestones, systems and owners.
2. Define what the agent may answer, recommend and execute.
3. Create the Copilot Studio agent purpose and instructions.
4. Add approved knowledge and a new-joiner profile input.
5. Design tools for task creation, reminder and status retrieval.
6. Add human confirmation before closing an access or equipment task.
7. Test normal joiner, missing manager, late equipment and confidential-question scenarios.
8. Review analytics and hand the pilot to named owners.

Prompt scaffold

You are supporting the Agentic AI for HR case. Use only the supplied synthetic evidence. Task: Design a Copilot Studio agent that coordinates onboarding while retaining accountable task owners. State assumptions, cite the evidence used, distinguish observation from inference, identify uncertainty, and list the human checks required before use.

Evidence to submit

- Conversation map
- Tool and approval design
- Test evidence

Acceptance criteria

- Least-privilege connectors
- No invented commitments
- Human owner confirms every completion

The submission is complete when the workbook evidence log contains the input/source, prompt or design decision, output, verification, revision and named human owner.

Scenario questions

26. Which step is orchestration rather than content generation?
27. Where could duplicate actions occur?
28. What evidence proves the joiner—not only the workflow—was ready?

Activity 9: Build a Leave Request Agent and Approval Flow

Scenario

Employees submit leave through messages that HR rekeys into a tracker. Errors occur around balance, overlapping dates and manager approvals.

Learning goal

Design a Copilot Studio agent plus deterministic flow for eligibility, approval, update and notification.

Tools

Microsoft Copilot Studio; Agent flows or Power Automate; Teams; Excel or HRIS sandbox

Procedure

1. Review the synthetic leave policy, balances and requests.
2. Define required inputs and validation messages.
3. Design the conversation to clarify dates, leave type and coverage.
4. Create a deterministic flow for balance check, overlap check and approval routing.
5. Add a confirmation step before the agent calls the flow.
6. Write successful, rejected and failed outcomes to the audit log.
7. Test insufficient balance, duplicate request, missing approver and connector failure.
8. Review access permissions and the employee-facing explanation.

Prompt scaffold

You are supporting the Agentic AI for HR case. Use only the supplied synthetic evidence. Task: Design a Copilot Studio agent plus deterministic flow for eligibility, approval, update and notification. State assumptions, cite the evidence used, distinguish observation from inference, identify uncertainty, and list the human checks required before use.

Evidence to submit

- Agent instructions
- Flow design
- Approval matrix
- Audit log and test report

Acceptance criteria

- Confirm before submission
- Manager approves exceptions
- Idempotency prevents duplicates

The submission is complete when the workbook evidence log contains the input/source, prompt or design decision, output, verification, revision and named human owner.

Scenario questions

29. Which rule belongs in the flow rather than the language model?
30. How will a duplicate submission be detected?
31. Who resolves a disputed balance?

Learning Unit 4: Responsible Adoption and HR Transformation

Govern decisions, data, people and agent operations

Concept notes

Why HR AI is high consequence

Employment decisions affect opportunity, income, dignity and trust.

- Recruitment and promotion shape opportunity
- Performance evidence can be incomplete
- Employee data is sensitive
- Power asymmetry raises consent concerns

Sources: <https://www.pdpc.gov.sg/help-and-resources/2020/01/model-ai-governance-framework>;
<https://www.mom.gov.sg/employment-practices/fair-consideration-framework>

Human involvement by impact

More consequential and less reversible decisions need stronger human authority.

- Inform: AI drafts or retrieves
- Recommend: AI proposes with evidence
- Approve: authorised human decides
- Execute: agent acts only within approved bounds

Sources: <https://www.pdpc.gov.sg/help-and-resources/2020/01/model-ai-governance-framework>;
<https://www.strategyand.pwc.com/de/en/functions/organisational-strategy/agentic-ai-in-hr.html>

Privacy by design

Minimise data before asking a model to protect it.

- Purpose and lawful use
- Minimum necessary fields
- Access, retention and deletion
- Incident and data-subject response

Sources: <https://www.pdpc.gov.sg/help-and-resources/2020/01/model-ai-governance-framework>

Fair hiring controls

Selection criteria must remain job-related, reviewable and consistently applied.

- Fairly consider candidates
- Exclude non-job-related protected traits
- Keep interview and decision evidence
- Provide human challenge and correction

Sources: <https://www.mom.gov.sg/employment-practices/fair-consideration-framework>;
<https://www.pdpc.gov.sg/help-and-resources/2020/01/model-ai-governance-framework>

Bias is a system property

Bias can enter through objectives, data, prompts, tools and human use.

- Question proxy variables
- Compare outcomes across relevant groups

- Review false positives and negatives
- Change process, not only the prompt

Sources: <https://www.pdpc.gov.sg/help-and-resources/2020/01/model-ai-governance-framework>;
<https://www.sap.com/resources/ai-for-hr>

Transparency and explainability

Employees need to know when AI is used and how to seek help.

- Disclose the AI role
- Show source and material factors
- Explain limits and uncertainty
- Provide an accessible human channel

Sources: <https://www.pdpc.gov.sg/help-and-resources/2020/01/model-ai-governance-framework>

Security and prompt injection

An agent can be manipulated through the text it reads.

- Treat retrieved content as untrusted
- Separate instructions from data
- Allow-list tools and destinations
- Monitor anomalous actions

Sources: <https://www.pdpc.gov.sg/help-and-resources/2020/01/model-ai-governance-framework>;
<https://www.microsoft.com/en-sg/microsoft-365-copilot/microsoft-copilot-studio>

Agent identity and least privilege

Every action should be attributable to an authorised identity.

- Use named service or user context
- Grant only required permissions
- Separate maker, approver and operator roles
- Review access and dormant agents

Sources: <https://www.microsoft.com/en-sg/microsoft-365-copilot/microsoft-copilot-studio>;
<https://www.ibm.com/think/topics/ai-agents-in-human-resources>

Operational controls

Agents are digital workers that need owners, logs, limits and retirement.

- Named business and technical owners
- Versioned instructions and knowledge
- Run-time monitoring and kill switch
- Periodic validation and decommissioning

Sources: <https://www.hrci.org/community/blogs-and-announcements/hr-leads-business-blog/hr-leads-business/2025/08/21/agentic-ai-in-hr>; <https://www.strategyand.pwc.com/de/en/functions/organisational-strategy/agentic-ai-in-hr.html>

Success measures

Measure service, adoption, quality, risk and human value together.

- Cycle time and first-contact resolution
- Accuracy and evidence completeness
- Escalation and override rates
- Employee trust and time returned to people

Sources: <https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/the-organization-blog/four-ways-to-start-using-generative-ai-in-hr>;
<https://www.gartner.com/en/human-resources/topics/artificial-intelligence-in-hr>

Change and AI literacy

Adoption depends on role clarity and practical learning.

- Explain what changes and what does not
- Train prompting, review and escalation
- Create champions and feedback loops
- Recognise new work: agent owner, tester, steward

Sources: <https://www.gartner.com/en/human-resources/topics/artificial-intelligence-in-hr>;
<https://www.auxis.com/ai-in-hr-examples-types-use-cases-best-practices/>

90-day responsible roadmap

Move from a bounded pilot to an owned capability through evidence gates.

- Days 1–30: baseline, risk and prototype
- Days 31–60: test, controls and user feedback
- Days 61–90: measured pilot and operating handover
- Scale only when value and control thresholds pass

Sources: <https://www.strategyand.pwc.com/de/en/functions/organisational-strategy/agentic-ai-in-hr.html>;
<https://www.auxis.com/ai-in-hr-examples-types-use-cases-best-practices/>

Activity 10: Prepare a Fair Performance and Learning Conversation

Scenario

A manager has fragmented notes, goals and service metrics for two employees and wants Copilot to draft development conversations.

Learning goal

Use Copilot to organise evidence, distinguish observation from inference and prepare human-led coaching.

Tools

Word and Teams with Microsoft 365 Copilot; Excel

Procedure

1. Review the synthetic goals, evidence and contextual notes.
2. Label each item as observation, result, feedback, context or unsupported inference.
3. Ask Copilot to summarise evidence by agreed goal and cite the source.
4. Challenge the summary for recency and confirmation bias.
5. Draft coaching questions that invite the employee's perspective.
6. Map agreed gaps to learning and stretch options.
7. Record employee corrections and final human judgement.
8. Check that the plan supports development rather than surveillance.

Prompt scaffold

You are supporting the Agentic AI for HR case. Use only the supplied synthetic evidence. Task: Use Copilot to organise evidence, distinguish observation from inference and prepare human-led coaching. State assumptions, cite the evidence used, distinguish observation from inference, identify uncertainty, and list the human checks required before use.

Evidence to submit

- Evidence summary
- Coaching questions
- Learning plan

Acceptance criteria

- No personality inference
- Employee can correct evidence
- Manager owns the conversation and rating

The submission is complete when the workbook evidence log contains the input/source, prompt or design decision, output, verification, revision and named human owner.

Scenario questions

32. Which evidence is relevant but not sufficient?
33. How can the employee contest the summary?
34. What should Copilot never infer from tone or activity data?

Activity 11: Design Safe Employee Service and Offboarding

Scenario

HR wants one conversational front door for benefits, payroll status, employee-relations triage and offboarding tasks.

Learning goal

Define privacy-aware routing, escalation and least-privilege actions across sensitive HR services.

Tools

Microsoft Copilot Studio; Teams; SharePoint; HRIS sandbox

Procedure

1. Classify the supplied intents by sensitivity, consequence and reversibility.
2. Decide which intents are answer, retrieve, transact or escalate.
3. Map authentication and minimum data for each service.
4. Design response patterns for payroll dispute, wellbeing disclosure and employee-relations complaint.
5. Build the offboarding sequence for confirmation, assets, access and knowledge transfer.
6. Assign maker, approver and verifier roles.
7. Test data leakage, social engineering, urgent safety and incomplete exit scenarios.
8. Record operational owner, retention and incident response.

Prompt scaffold

You are supporting the Agentic AI for HR case. Use only the supplied synthetic evidence. Task: Define privacy-aware routing, escalation and least-privilege actions across sensitive HR services. State assumptions, cite the evidence used, distinguish observation from inference, identify uncertainty, and list the human checks required before use.

Evidence to submit

- Intent-and-risk catalogue
- Escalation matrix
- Offboarding control map

Acceptance criteria

- Authenticate personal requests
- No legal or clinical advice
- Separation of duties for access removal

The submission is complete when the workbook evidence log contains the input/source, prompt or design decision, output, verification, revision and named human owner.

Scenario questions

35. Why should one front door not mean one permission set?
36. Which disclosure demands an immediate human response?
37. What evidence is needed before access is removed?

Activity 12: Capstone: Operate the HR Agent Control Room

Scenario

HarbourLight's pilot portfolio now includes policy, onboarding and leave agents. Leaders must decide whether to scale, remediate or stop each agent.

Learning goal

Use service, risk, adoption and human-value evidence to make a defensible operating decision.

Tools

Excel; Teams; Microsoft 365 Copilot

Procedure

1. Review the synthetic KPI, incident, test and feedback dataset.
2. Validate denominator, target and trend for each metric.
3. Ask Copilot to summarise evidence but require cell references.
4. Identify control failures, unintended outcomes and data gaps.
5. Classify each agent as scale, remediate, pause or retire.
6. Respond to the prompt-injection incident using the provided playbook.
7. Create a 90-day roadmap with decision gates and accountable owners.
8. Present a three-minute executive recommendation and answer peer challenge questions.

Prompt scaffold

You are supporting the Agentic AI for HR case. Use only the supplied synthetic evidence. Task: Use service, risk, adoption and human-value evidence to make a defensible operating decision. State assumptions, cite the evidence used, distinguish observation from inference, identify uncertainty, and list the human checks required before use.

Evidence to submit

- Agent scorecard
- Incident response
- 90-day roadmap and executive recommendation

Acceptance criteria

- No scale without control thresholds
- Employee feedback included
- Named owner and retirement path

The submission is complete when the workbook evidence log contains the input/source, prompt or design decision, output, verification, revision and named human owner.

Scenario questions

38. Which metric looked good but hid harm?
39. What evidence justifies scale?
40. How will the organisation retire an agent safely?

Responsible HR AI checklist

- ☐ Purpose and scope are explicit.
- ☐ Only necessary and authorised data are used.
- ☐ Job-related evidence and fair criteria are documented.
- ☐ Sources, uncertainty and limitations are visible.
- ☐ Consequential decisions have an authorised human approver.
- ☐ Agent identity, permissions, tools and destinations use least privilege.
- ☐ Normal, edge, adversarial and failure cases are tested.
- ☐ Logs, incidents, feedback, change and retirement have owners.

Source register

The complete reviewed source register, access status and cross-artifact coverage matrix are maintained in SOURCE-COVERAGE.md alongside the courseware. Key authoritative anchors include Microsoft product/adoption guidance, Singapore PDPC's Model AI Governance Framework and Singapore MOM's Fair Consideration Framework.