

TERTIARY INFOTECH ACADEMY

LESSON PLAN

Agentic AI for HR

Course Code: TGS-2024045795

2 days / 16 hours

Human Resource Digitalisation-4 (HRS-HRM-4031-1.1)

Version 6.0 • 16 August 2026

Course and delivery overview

Field	Detail
Course	Agentic AI for HR
Code	TGS-2024045795
Duration	2 days, 09:00–18:00 daily; 16 training hours excluding lunch
Mode	Facilitated classroom or synchronous virtual delivery
Assessment	Written Assessment, 60 minutes; Case Study Assessment, 60 minutes

Learning outcomes and evidence

LO1: Compare how generative and agentic AI reshape HR operating models versus conventional HR delivery.

LO2: Analyse where Microsoft 365 Copilot and AI agents can improve end-to-end HR processes and employee experience.

LO3: Conduct feasibility analysis and design a governed implementation for an HR generative-AI or agentic-AI use case.

LO4: Apply legal, ethical, privacy, security, fairness and human-oversight controls to AI-enabled HR work.

Facilitation principles

- Use the slides for visual explanation; use the Learner Guide for procedures.
- Demonstrate only with synthetic records and approved tenant content.
- Ask learners to state evidence, uncertainty and human ownership.
- Do not coach or disclose answers during individual assessment.

Day 1 schedule

Time	Topic/activity	Method	Slide/evidence alignment
09:00–09:30	Welcome, outcomes and assessment briefing	Facilitated briefing	Slides 1–15
09:30–10:30	Unit 1 concepts	Discussion and visual examples	Slides 17–36
10:30–10:45	Break	—	—
10:45–12:30	Activities 1–3	Team practice and review	A1: 30–31; A2: 32–33; A3: 34–35
12:30–13:30	Lunch	—	—
13:30–15:00	Unit 2 concepts	Demonstration and discussion	Slides 37–56
15:00–15:15	Break	—	—
15:15–17:30	Activities 4–6	Individual and team practice	A4: 50–51; A5: 52–53; A6: 54–55
17:30–18:00	Day 1 integration and reflection	Peer teach-back	Recap

Trainer notes

Establish the continuum from deterministic automation to Copilot-assisted work and bounded agency. Challenge learners to choose the lightest sufficient capability. During recruitment and policy activities, repeatedly separate evidence from inference and require human approval.

Day 2 schedule

Time	Topic/activity	Method	Slide/evidence alignment
09:00–09:15	Day 1 retrieval practice	Facilitated review	Recap
09:15–10:30	Unit 3 concepts	Agent design walkthrough	Slides 57–76
10:30–10:45	Break	—	—
10:45–12:30	Activities 7–9	Build, test and review	A7: 70–71; A8: 72–73; A9: 74–75
12:30–13:30	Lunch	—	—
13:30–14:30	Unit 4 concepts	Governance scenario analysis	Slides 77–96
14:30–14:45	Break	—	—
14:45–15:30	Activities 10–12 capstone	Evidence-based decision	A10: 90–91; A11: 92–93; A12: 94–95
15:30–16:00	Assessment preparation	Questions and readiness	Assessment briefing
16:00–17:00	Written Assessment	Individual assessment	K1–K5
17:00–18:00	Case Study Assessment	Individual assessment	A1–A5

Trainer notes

Keep agent building anchored to purpose, knowledge, tools, boundaries and tests. Rehearse failure and escalation. Before assessment, clarify logistics and permitted resources without revealing model responses.

Activity facilitation and debrief

Activity 1: Map the HR AI Opportunity Portfolio

Purpose: Classify HR work across automation, Microsoft 365 Copilot and agentic AI, then select one bounded pilot.

Observe for: No employee personal data; State all assumptions; Human owner for every shortlisted use case

Debrief: Why is the selected capability the lightest sufficient option? | What evidence would stop the pilot? | Whose experience could worsen even if cycle time improves?

Slides: 30–31

Activity 2: Design a Skills-Based Workforce Plan

Purpose: Use Copilot to synthesise demand evidence and draft a skills-based, fair job design.

Observe for: Separate evidence from assumption; No age, nationality or other non-job-related criteria; Hiring manager approval

Debrief: Which assumption has the greatest planning impact? | Which requirement could unfairly narrow the pool? | How will the hypothesis be revisited?

Slides: 32–33

Activity 3: Write the Governed HR AI Pilot Charter

Purpose: Convert an opportunity into a testable pilot charter with go/no-go gates.

Observe for: Synthetic or de-identified pilot data; Named risk owner; Kill switch and rollback

Debrief: Which decision remains human and why? | Who can stop the pilot? | What would a misleading success metric reward?

Slides: 34–35

Activity 4: Build a Fair Resume Evidence Matrix

Purpose: Use Microsoft 365 Copilot to extract job-related evidence and prepare a human-reviewed shortlist.

Observe for: Synthetic resumes only; Protected attributes excluded; AI cannot reject candidates

Debrief: Where did missing evidence risk becoming a negative assumption? | What would make this screening process contestable? | Which errors matter most: false inclusion or false exclusion?

Slides: 50–51

Activity 5: Draft and Govern an HR Policy Update

Purpose: Use Word Copilot to compare sources and prepare a controlled policy draft plus communication plan.

Observe for: Only approved source documents; Policy owner signs off; Effective date and superseded versions recorded

Debrief: Which conflict could create unequal treatment? | What must not be inferred by Copilot? | How will users know which version is authoritative?

Slides: 52–53

Activity 6: Create a Personalised Onboarding Journey

Purpose: Use Microsoft 365 Copilot to prepare role-based onboarding plans, communications and learning prompts.

Observe for: Use only approved role and policy data; Do not infer personal needs; New joiner can request adjustments

Debrief: What should remain standard for every new joiner? | Which personalisation requires consent? | How will HR detect an onboarding access failure early?

Slides: 54–55

Activity 7: Build an HR Policy Q&A Agent

Purpose: Create a bounded policy agent in Agent Builder for Microsoft 365 Copilot and test grounded answers.

Observe for: Cite approved policy; No personalised entitlement decisions; Escalate ambiguity and sensitive cases

Debrief: What makes an answer grounded rather than merely plausible? | When should the agent stop answering? | What new operating work does this agent create?

Slides: 70–71

Activity 8: Design an Onboarding Coordination Agent

Purpose: Design a Copilot Studio agent that coordinates onboarding while retaining accountable task owners.

Observe for: Least-privilege connectors; No invented commitments; Human owner confirms every completion

Debrief: Which step is orchestration rather than content generation? | Where could duplicate actions occur? | What evidence proves the joiner—not only the workflow—was ready?

Slides: 72–73

Activity 9: Build a Leave Request Agent and Approval Flow

Purpose: Design a Copilot Studio agent plus deterministic flow for eligibility, approval, update and notification.

Observe for: Confirm before submission; Manager approves exceptions; Idempotency prevents duplicates

Debrief: Which rule belongs in the flow rather than the language model? | How will a duplicate submission be detected? | Who resolves a disputed balance?

Slides: 74–75

Activity 10: Prepare a Fair Performance and Learning Conversation

Purpose: Use Copilot to organise evidence, distinguish observation from inference and prepare human-led coaching.

Observe for: No personality inference; Employee can correct evidence; Manager owns the conversation and rating

Debrief: Which evidence is relevant but not sufficient? | How can the employee contest the summary? | What should Copilot never infer from tone or activity data?

Slides: 90–91

Activity 11: Design Safe Employee Service and Offboarding

Purpose: Define privacy-aware routing, escalation and least-privilege actions across sensitive HR services.

Observe for: Authenticate personal requests; No legal or clinical advice; Separation of duties for access removal

Debrief: Why should one front door not mean one permission set? | Which disclosure demands an immediate human response? | What evidence is needed before access is removed?

Slides: 92–93

Activity 12: Capstone: Operate the HR Agent Control Room

Purpose: Use service, risk, adoption and human-value evidence to make a defensible operating decision.

Observe for: No scale without control thresholds; Employee feedback included; Named owner and retirement path

Debrief: Which metric looked good but hid harm? | What evidence justifies scale? | How will the organisation retire an agent safely?

Slides: 94–95