



TERTIARY INFOTECH ACADEMY • WSQ

LESSON PLAN

Microsoft Copilot for HR

Course Code: TGS-2024045795

2 days / 16 hours

Human Resource Digitalisation-4 (HRS-HRM-4031-1.1)

UEN 201200696W

Version 6.3 • 27 September 2026

Document Version Control

Version	Date	Summary
6.0	16 August 2026	Initial redesigned courseware package
6.1	16 August 2026	SPC-aligned administration pages and visual repairs
6.2	16 August 2026	Expanded research, full HR lifecycle activities, mock data and learner PDFs
6.3	27 September 2026	Renamed Microsoft Copilot for HR; converted activities to numbered labs and aligned the assessment

Version numbers apply to this generated courseware release; assessment instruments retain their separately controlled version.

Contents

Course and delivery overview

Field	Detail
Course	Microsoft Copilot for HR
Code	TGS-2024045795
Duration	2 days, 09:30–18:30 daily; 16 scheduled contact hours excluding lunch, including two facilitated 15-minute reflection intervals per day
Mode	Facilitated classroom or synchronous virtual delivery
Assessment	Written Assessment, 60 minutes; Case Study Assessment, 60 minutes

Learning outcomes and evidence

LO1: Compare how generative and agentic AI reshape HR operating models versus conventional HR delivery.

LO2: Analyse where Microsoft 365 Copilot and AI agents can improve end-to-end HR processes and employee experience.

LO3: Conduct feasibility analysis and design a governed implementation for an HR generative-AI or agentic-AI use case.

LO4: Apply legal, ethical, privacy, security, fairness and human-oversight controls to AI-enabled HR work.

Labs 1–3, 4–6 and 7–9 are parallel team pods: each learner completes one lab in class and shares the result with peers; the other two are self-paced. Labs 10–11 are self-paced extensions. Lab 12 is a 45-minute in-class decision sprint using supplied evidence; full build and regression testing continue after class. The individual WA and Case Study remain 60 minutes each.

Facilitation principles

- Open with the current HR pain points and return throughout to the three-part value message: less administration, higher productivity and more strategic HR capacity.
- Use the slides for visual explanation; use the Learner Guide for procedures.
- Demonstrate only with synthetic records and approved tenant content.
- Ask learners to state evidence, uncertainty and human ownership.
- Do not coach or disclose answers during individual assessment.

Day 1 schedule

Time	Topic/activity	Method	Slide/evidence alignment
09:30–10:00	Welcome, outcomes and assessment briefing	Facilitated briefing	Slides 1–16
10:00–11:00	Unit 1 concepts	Discussion and visual examples	Slides 17–63
11:00–11:15	Facilitated reflection	Prompt check and peer debrief	Current unit/lab
11:15–13:00	Labs 1–3	Team practice and review	L1: 36–44; L2: 45–53; L3: 54–62
13:00–14:00	Lunch	—	—
14:00–15:30	Unit 2 concepts	Demonstration and discussion	Slides 64–106
15:30–15:45	Facilitated reflection	Prompt check and peer debrief	Current unit/lab
15:45–18:00	Labs 4–6	Individual and team practice	L4: 79–87; L5: 88–96; L6: 97–105
18:00–18:30	Day 1 integration and reflection	Peer teach-back	Recap

Trainer notes

Begin with fragmented requests, repeated rekeying, slow responses and weak traceability. Establish the continuum from deterministic automation to Copilot-assisted work and bounded agency. Use Labs 1–6 as one talent journey from demand through development. At every debrief, ask what administration was removed, what productivity improved and what strategic or human work received the returned capacity.

Day 2 schedule

Time	Topic/activity	Method	Slide/evidence alignment
09:30–09:45	Day 1 retrieval practice	Facilitated review	Recap
09:45–11:00	Unit 3 concepts	Agent design walkthrough	Slides 107–148
11:00–11:15	Facilitated reflection	Prompt check and peer debrief	Current unit/lab
11:15–13:00	Labs 7–9	Build, test and review	L7: 121–129; L8: 130–138; L9: 139–147
13:00–14:00	Lunch	—	—
14:00–15:00	Unit 4 concepts	Governance scenario analysis	Slides 149–193
15:00–15:15	Facilitated reflection	Prompt check and peer debrief	Current unit/lab
15:15–16:00	Labs 10–12 capstone	Evidence-based decision	L10: 166–174; L11: 175–183; L12: 184–192
16:00–16:30	Assessment preparation	Questions and readiness	Assessment briefing
16:30–17:30	Written Assessment	Individual assessment	K1–K5
17:30–18:30	Case Study Assessment	Individual assessment	A1–A5

Trainer notes

Use Labs 7–9 as parallel employee-service pods and Lab 12 as a short control-room decision sprint; Labs 10–11 are self-paced extensions. Use the employee-service and exit journey. Keep agent building anchored to purpose, knowledge, tools, boundaries and tests. Rehearse failure and escalation, and distinguish time genuinely returned from new monitoring work. Before assessment, clarify logistics and permitted resources without revealing model responses.

Lab facilitation and debrief

Lab 1: Assess Talent Needs with Copilot

Purpose: Use Microsoft 365 Copilot to turn operational evidence into a transparent hire, redeploy or develop recommendation.

Observe for: Cite the supplied cells and records; Separate fact from forecast; Business and HR owners approve the demand hypothesis

Debrief: Which assumption most changes the talent need? | Where could automation reduce workload without removing accountability? | What evidence would trigger a plan revision?

Slides: 36–44

Lab 2: Draft an Inclusive Job Description with Generative AI

Purpose: Use Word Copilot to draft a skills-based, inclusive and approval-ready job description grounded in the approved talent need.

Observe for: Use job-related criteria only; Distinguish essential from trainable skills; Hiring manager and HR approve every requirement

Debrief: Which copied requirement would unnecessarily narrow the talent pool? | What must the hiring manager verify rather than delegate to Copilot? | How does the JD make later screening contestable?

Slides: 45–53

Lab 3: Screen Resumes with a Fair Evidence Matrix

Purpose: Use Microsoft 365 Copilot to extract job-related evidence at speed while keeping shortlist decisions human, traceable and reviewable.

Observe for: Synthetic resumes only; Protected attributes and proxies excluded; AI may organise evidence but cannot reject a candidate

Debrief: Where could absence of evidence become an unfair negative assumption? | How can a candidate challenge or correct the evidence? | Which screening error creates the greatest harm?

Slides: 54–62

Lab 4: Prepare and Review Structured Interviews with AI

Purpose: Use Copilot to prepare job-related structured questions, organise synthetic interview notes and support a human panel decision.

Observe for: Same core questions for every candidate; No emotion or personality inference; Panel members score independently before consensus

Debrief: Which question best tests a role outcome rather than similarity to the panel? | What interview evidence should Copilot never infer? | How will disagreements between panel members be resolved?

Slides: 79–87

Lab 5: Automate Onboarding Coordination

Purpose: Design a Copilot Studio onboarding agent that coordinates routine work while retaining accountable task owners and human confirmation.

Observe for: Use approved role and policy data; Least-privilege connectors; A human owner confirms every completion

Debrief: Which work is safely automated and which remains owned by a person? | Where could duplicate tasks be created? | What proves the new joiner is ready rather than the checklist merely complete?

Slides: 88–96

Lab 6: Create a Personalised Talent Development Journey

Purpose: Use Copilot and a bounded development agent to identify evidence-based skill gaps, prepare learning options and automate routine follow-up.

Observe for: Employee can correct the skills record; No deterministic potential ranking; Manager and employee choose the final plan

Debrief: How can the employee challenge the skills profile? | What would make a recommendation explainable? | Which development decision should never be automated?

Slides: 97–105

Lab 7: Draft and Operate an HR Policy Agent

Purpose: Use Word Copilot to reconcile a controlled policy draft, then build an Agent Builder policy Q&A agent that cites the approved version.

Observe for: Policy owner approves meaning; Answers cite effective policy; Ambiguous and sensitive cases escalate

Debrief: What makes an answer grounded rather than plausible? | Which exception requires a policy owner instead of an agent? | What operating work is created when self-service is introduced?

Slides: 121–129

Lab 8: Automate Leave Application and Approval

Purpose: Build a Copilot Studio leave agent plus deterministic approval flow for validation, confirmation, update and notification.

Observe for: Employee confirms before submission; Manager owns exceptions; Idempotency prevents duplicate actions

Debrief: Which rule belongs in the flow rather than the language model? | How will the system detect a duplicate request? | Who resolves a disputed leave balance?

Slides: 130–138

Lab 9: Build a Benefits Self-Service Agent

Purpose: Design a benefits agent that answers grounded general questions, retrieves authorised status and escalates consequential or sensitive cases.

Observe for: General answers cite approved plan documents; Personal requests require authentication; No medical, legal or final entitlement advice

Debrief: Why should one conversational front door not have one permission set? | Which benefits question requires immediate human review? | How will employees know the source and effective date of an answer?

Slides: 139–147

Lab 10: Use AI for Workforce Wellness Insights

Purpose: Use Copilot to analyse aggregated wellness evidence, identify organisational questions and prepare human-led interventions.

Observe for: Aggregate and minimise personal data; No diagnosis or individual risk scoring; Employees have a clear human support route

Debrief: Which attractive insight is not safe to act on? | How does aggregation reduce but not eliminate privacy risk? | What response supports people without surveillance?

Slides: 166–174

Lab 11: Automate Safe Employee Offboarding

Purpose: Design a Copilot Studio offboarding agent and controlled flow that reduces follow-up while preserving separation of duties and respectful communication.

Observe for: Confirmed exit event before action; Separate HR, payroll, asset and access approvals; Knowledge and personal data retention are purpose-limited

Debrief: What evidence is required before access is removed? | Where is separation of duties essential? | How can automation preserve dignity during exit?

Slides: 175–183

Lab 12: Capstone: Operate the HR Agent Control Room

Purpose: Use lifecycle, service, risk, adoption and human-value evidence to govern the full HR AI operating system.

Observe for: No scale without control thresholds; Employee feedback included; Named owner and retirement path

Debrief: Which metric looked good but hid harm? | What evidence justifies scale? | How will the organisation retire an agent safely?

Slides: 184–192